

ADL Hw3

R11944014 戴靖婷

tags: ADL Python Report

Q1: Model

Model: Describe the model architecture and how it works on text summarization.

- **google/mt5-small:**
 - Multilingual T5 (mT5) is a text-to-text, encoder-decoder transformer with attention mechanism, which means it figures out the relationship between all words in a sentence.
 - Also, it is pretrained on multilingual Common Crawl corpus (mC4).
 - The following is the content of config.json.

```
{
  "_name_or_path": "google/mt5-small",
  "architectures": [
    "MT5ForConditionalGeneration"
  ],
  "d_ff": 1024,
  "d_kv": 64,
  "d_model": 512,
  "decoder_start_token_id": 0,
  "dense_act_fn": "gelu_new",
  "dropout_rate": 0.1,
  "eos_token_id": 1,
  "feed_forward_proj": "gated-gelu",
  "initializer_factor": 1.0,
  "is_encoder_decoder": true,
  "is_gated_act": true,
  "layer_norm_epsilon": 1e-06,
  "model_type": "mt5",
  "num_decoder_layers": 8,
  "num_heads": 6,
  "num_layers": 8,
  "pad_token_id": 0,
  "relative_attention_max_distance": 128,
  "relative_attention_num_buckets": 32,
  "tie_word_embeddings": false,
  "tokenizer_class": "T5Tokenizer",
  "torch_dtype": "float32",
  "transformers_version": "4.22.2",
  "use_cache": true,
  "vocab_size": 250100
}
```

- For text summarization,

- It can be seen as a machine translation task by translating a text into a shorter summarization.
- To specify the task of mt5, add a prompt prefix `--source_prefix "summarize: "`.
- I made use of the sameple code from the huggingface github.
<https://github.com/huggingface/transformers/tree/t5-fp16-no-nans/examples/pytorch/summarization> (<https://github.com/huggingface/transformers/tree/t5-fp16-no-nans/examples/pytorch/summarization>)

Preprocessing: Describe your preprocessing (e.g. tokenization, data cleaning and etc.)

- Tokenization:
 - Tokenizer class: T5Tokenizer
 - T5Tokenizer is based on SentencePiece method.
 - It treats the input as the raw input including space because not every language separate sentences by a space, such as Chinese.
 - Then, it constructs vocabulary by the BPE or unigram tokenization algorithms.
 - Therefore, it is suitable for multilingual T5 and also builds an independent system for each language end-to-end.


Q2: Training

Hyperparameter: Describe your hyperparameter you use and how you decide it.

```
python ./train_summarization.py \
  --model_name_or_path google/mt5-small \
  --do_train \
  --do_eval \
  --train_file ./data/train.jsonl \
  --validation_file ./data/public.jsonl \
  --source_prefix "summarize: " \
  --text_column maintext \
  --summary_column title \
  --per_device_train_batch_size 4 \
  --per_device_eval_batch_size 4 \
  --optim adafactor \
  --learning_rate 5e-4 \
  --max_source_length 256 \
  --max_target_length 64 \
  --output_dir ./ckpt/ \
  --overwrite_output_dir \
  --predict_with_generate \
  --evaluation_strategy steps \
  --eval_steps 500 \
  --logging_steps 500
```

- I trained my model with 1024 and 256 at first. However, it used more than 8G memory, so I truncated the text length.

Learning Curves: Plot the learning curves (ROUGE versus training steps)



- This result is generated by greedy method with fp32.

Q3: Generation Strategies

Strategies: Describe the detail of the following generation strategies:

- Greedy
 - For the next word, select the word with the highest probability.
 - $w_t = \text{argmax}_w P(w|w_{1:t-1})$
- Beam Search
 - Keep the most likely num_beams of possible words at each time and choose the one with overall highest probability.
 - It avoids missing the hidden high probability word sequences behind a relatively lower probability word.
- Top-k Sampling
 - Keep only the K most likely next words for sampling from the reconstructed distribution.
 - And, sampling means picking the next word based on its conditional probability distribution. $w_t \sim P(w|w_{1:t-1})$
- Top-p Sampling
 - Choose the next word from the set of words whose cumulative probability exceeds the probability p.
- Temperature
 - Temperature changes the probability of every word.

- The range of temperature is between 0.0 to 1.0.
 - When temperature is closer to 0.0, it is closer to greedy.
 - When temperature is closer to 1.0, each word has the same probability to be generated as the next word.
- Lower the temperature of the softmax makes the distribution P sharper, which means increasing the likelihood of the words with high probability and decreasing the likelihood of the words with low probability.
- Reference: <https://huggingface.co/blog/how-to-generate> (<https://huggingface.co/blog/how-to-generate>)

Hyperparameters

- Try at least 2 settings of each strategies and compare the result. The ROUGE result is evaluated by `--do_eval` in `run_summarization.py` instead of `eval.py` (<http://eval.py>).
 - Greedy
 - The more accurate the floating point number (fp32) is, the better the result is.

	rouge-1	rouge-2	rouge-l
fp32	27.1872	10.1801	24.2615
fp16	26.9455	9.981	23.9738

- Beam Search
 - Adjust `num_beams` and set `early_stopping=True`.
 - The larger the `num_beams` is, the more possible predictions it keeps. So the result is better when the `num_beams` is larger.

	rouge-1	rouge-2	rouge-l
<code>num_beams = 2</code>	28.0982	11.0138	25.0553
<code>num_beams = 5</code>	28.3986	11.4707	25.3744
<code>num_beams = 10</code>	28.3385	11.5820	25.2885

- Top-k Sampling
 - Adjust `top_k` and set `do_sample=True`.
 - The result is not good if using a large set of likely next words, so we may narrow down the set of the likely next words.

	rouge-1	rouge-2	rouge-l
top_k = 20	23.0773	7.2718	20.1300
top_k = 10	24.2822	8.0891	21.3037
top_k = 5	25.5530	8.7149	22.3640

- Top-p Sampling

- Adjust top_p and set top_k=0, do_sample=True.
- The result is better when p is smaller, which means using less likely word set is better. The trend of top_p is the same as the trend of top_k.

-

	rouge-1	rouge-2	rouge-l
top_p = 0.95	18.7501	5.6598	16.5373
top_p = 0.50	24.2476	8.5073	21.3852
top_p = 0.30	25.7993	9.2286	22.7769

- Temperature

- Adjust temperature and set top_k=0, do_sample=True.
- It is better to set temperature closer to 0, it is closer to greedy.

-

	rouge-1	rouge-2	rouge-l
temperature = 0.7	24.1415	8.4414	21.2756
temperature = 0.5	25.9870	9.4473	22.9505
temperature = 0.3	26.9834	9.9318	24.0372

- What is your final generation strategy? (you can combine any of them)
 - Over the experiments, I found that the result of beam search is the best, so my final generation strategy is **num_beams=5** with early_stopping=True and the test file is predicted from it.